

2.6 LU and PLU Factorizations

Gaussian elimination solves one system $A\mathbf{x} = \mathbf{b}$. In practice the matrix often stays fixed while the right-hand side changes: the same network under different loads, the same model fed different data. The row operations depend only on A , so repeating them for every $\mathbf{b}$ wastes almost all of the work. This section stores the elimination as a factorization $A = LU$, a lower triangular matrix that remembers the operations times an upper triangular matrix that is their result. We prove when it exists and when it is unique, and we repair its failures with row swaps.

Throughout, F is a field and all matrices are square, in $M_n(F)$.

2.6.1 Triangular systems are cheap

Some systems need no elimination at all. Take

$$U = \begin{pmatrix} 2 & 1 & -1 \\ 0 & 3 & 1 \\ 0 & 0 & -2 \end{pmatrix}, \quad U\mathbf{x} = \begin{pmatrix} -1 \\ -1 \\ -4 \end{pmatrix}.$$

The last equation mentions only x_3 : $-2x_3 = -4$, so $x_3 = 2$. The second mentions x_2 and the known x_3 : $3x_2 + 2 = -1$, so $x_2 = -1$. The first then gives $2x_1 - 1 - 2 = -1$, so $x_1 = 1$. This is **back substitution**: in an upper triangular system with non-zero diagonal entries, each equation introduces exactly one new unknown, read from the bottom up. A lower triangular system is solved the same way from the top down, by **forward substitution**.

Now suppose a matrix A factors as $A = LU$ with L lower triangular and U upper triangular. Then $A\mathbf{x} = \mathbf{b}$ reads $L(U\mathbf{x}) = \mathbf{b}$. Naming the inner vector $\mathbf{y} = U\mathbf{x}$ splits the problem in two:

first solve $L\mathbf{y} = \mathbf{b}$ for $\mathbf{y}$, then solve $U\mathbf{x} = \mathbf{y}$ for $\mathbf{x}$.

Both are triangular, so both are cheap. The expensive part, finding L and U , is done once for all right-hand sides.

An LU factorization is Gaussian elimination written down once, so that it never has to be done again.

2.6.2 LU factorizations

Definition 2.6.1 — LU factorization

A matrix $L \in M_n(F)$ is **unit lower triangular** if it is lower triangular and **every** diagonal entry equals 1. Let $A \in M_n(F)$. An **LU factorization** of A is a pair (L, U) of matrices in $M_n(F)$ such that

$$A = LU, \quad L \text{ is **unit** lower triangular,} \quad U \text{ is upper triangular.}$$

In words: L has 1's on the diagonal and zeros above it, and the entries below the diagonal are unrestricted. U has zeros strictly below the diagonal, and its diagonal entries are unrestricted; some of them may be 0. The definition says **an** factorization: whether there is only one is a theorem, and it needs a hypothesis.

Example. (An LU factorization from elimination)

Let

$$A = \begin{pmatrix} 2 & 1 & -1 \\ 4 & 5 & -1 \\ -2 & 8 & 2 \end{pmatrix} \in M_3(\mathbb{Q}).$$

Reduce A to upper triangular form using only operations “subtract a multiple of a higher row from a lower row”, and record each multiplier. Compare the result with the matrix U of the previous subsection.

Solution. The first pivot is 2. To clear column 1 below it we subtract $\ell_{21} = 4/2 = 2$ times row 1 from row 2, and $\ell_{31} = -2/2 = -1$ times row 1 from row 3:

$$\begin{pmatrix} 2 & 1 & -1 \\ 4 & 5 & -1 \\ -2 & 8 & 2 \end{pmatrix} \xrightarrow[R_3 \rightarrow R_3 - (-1)R_1]{R_2 \rightarrow R_2 - 2R_1} \begin{pmatrix} 2 & 1 & -1 \\ 0 & 3 & 1 \\ 0 & 9 & 1 \end{pmatrix}.$$

The second pivot is 3. We subtract $\ell_{32} = 9/3 = 3$ times row 2 from row 3:

$$\begin{pmatrix} 2 & 1 & -1 \\ 0 & 3 & 1 \\ 0 & 9 & 1 \end{pmatrix} \xrightarrow{R_3 \rightarrow R_3 - 3R_2} \begin{pmatrix} 2 & 1 & -1 \\ 0 & 3 & 1 \\ 0 & 0 & -2 \end{pmatrix} = U.$$

Put the multipliers below the diagonal of a unit lower triangular matrix, each in the position of the entry it cleared:

$$L = \begin{pmatrix} 1 & 0 & 0 \\ \ell_{21} & 1 & 0 \\ \ell_{31} & \ell_{32} & 1 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ -1 & 3 & 1 \end{pmatrix}.$$

Multiplying out, row by row,

$$LU = \begin{pmatrix} 2 & 1 & -1 \\ 2 \cdot 2 & 2 \cdot 1 + 3 & 2 \cdot (-1) + 1 \\ -2 & -1 + 9 & 1 + 3 - 2 \end{pmatrix} = \begin{pmatrix} 2 & 1 & -1 \\ 4 & 5 & -1 \\ -2 & 8 & 2 \end{pmatrix} = A.$$

So (L, U) is an LU factorization of A , and U is exactly the triangular matrix of the back substitution above.

The example suggests the general statement: elimination produces U , and the multipliers, **with their signs as subtracted**, fill in L . We prove this below. First, the smallest cases.

- **Upper triangular A .** $A = I \cdot A$, and I is unit lower triangular. So every upper triangular matrix, including 0 and I , has the LU factorization (I, A) .
- 1×1 . Every $(a) \in M_1(F)$ is upper triangular, so $(a) = (1)(a)$. Nothing can go wrong here, and that is exactly why the 2×2 failure below is instructive.

Non-example by minimal change. The matrix $\begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix}$ has the LU factorization

$$\begin{pmatrix} 1 & 1 \\ 1 & 0 \end{pmatrix} = \begin{pmatrix} 1 & 0 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} 1 & 1 \\ 0 & -1 \end{pmatrix}.$$

Change the top-left entry from 1 to 0, and everything collapses.

Example. (An invertible matrix with no LU factorization)

Show that $A = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix} \in M_2(F)$ has **no** LU factorization, over any field F , although A is invertible.

Solution. A is invertible, since $AA = I_2$. Suppose, for a contradiction, that $A = LU$ with

$$L = \begin{pmatrix} 1 & 0 \\ \ell & 1 \end{pmatrix}, \quad U = \begin{pmatrix} u & v \\ 0 & w \end{pmatrix}.$$

Then

$$LU = \begin{pmatrix} u & v \\ \ell u & \ell v + w \end{pmatrix}.$$

Comparing $(1, 1)$ -entries with A gives $u = 0$. Comparing $(2, 1)$ -entries gives $\ell u = 1$. But $\ell u = \ell \cdot 0 = 0$ by Theorem 0.7.1, and $0 \neq 1$ in a field. This contradiction shows that no LU factorization exists.

What failed is visible: the $(1, 1)$ -entry of any product LU is u_{11} alone, so a zero there forces the whole first column of LU to vanish. In elimination terms, the first pivot is 0 while the entry below it is not, and the only cure is to **swap** rows, which “subtract a multiple of a higher row from a lower row” cannot do. We return to this with permutation matrices.

Why “unit”. Without the 1’s on the diagonal of L , uniqueness would fail even for invertible matrices. If $A = LU$ and D is any invertible diagonal matrix, then also $A = (LD)(D^{-1}U)$, where LD is still lower triangular and $D^{-1}U$ still upper triangular. When L is invertible and $D \neq I$, this is a different pair, and over any field with more than two elements such a D exists. For instance $(2) = (1) (2) = (2) (1) = (4) (1/2)$ over $\mathbb{Q}$. Fixing the diagonal of L removes this freedom and matches what elimination does: it never rescales a row.

Warning. **Invertible does not mean “has an LU factorization”.** The matrix of [Example \(An invertible matrix with no LU factorization\)](#) is invertible and has none. Conversely, a singular matrix can have one: $0 = I \cdot 0$. Existence of an LU factorization is a property of the **order of the rows**, not of invertibility.

2.6.3 Unit lower triangular matrices

Elimination builds L out of elementary matrices, and the uniqueness proof below divides one factorization by another. Both need the same facts: products and inverses of unit lower triangular matrices stay unit lower triangular. We also need to know exactly when a triangular matrix is invertible.

The key observation is a statement about where the non-zero entries of a product can sit.

Lemma 2.6.1 — Unit lower triangular matrices are closed

Let $L, L' \in M_n(F)$ be unit lower triangular. Then:

- (a) LL' is unit lower triangular;
- (b) L is invertible, and L^{-1} is unit lower triangular.

Idea. Write $L = I + N$, where N has non-zero entries only strictly below the diagonal. Multiplying two such matrices pushes the non-zero entries **further** below the diagonal, so $N^n = 0$. Then the geometric series $1/(1+t) = 1 - t + t^2 - \dots$ stops after n terms and gives the inverse exactly.

Proof. Say that $X \in M_n(F)$ is **p -low** if $x_{ij} = 0$ whenever $i < j + p$. So 0-low means lower triangular, and 1-low means zero on and above the diagonal.

Claim. If X is p -low and Y is q -low, then XY is $(p+q)$ -low. A matrix that is n -low is the zero matrix.

Proof. By Definition 0.9.9, $(XY)_{ij} = \sum_{l=1}^n x_{il}y_{lj}$. A term can be non-zero only if $i \geq l + p$ and $l \geq j + q$, which together give $i \geq j + p + q$. So if $i < j + p + q$, every term is 0 and $(XY)_{ij} = 0$. For the second statement, every pair of indices satisfies $i \leq n < 1 + n \leq j + n$, so every entry of an n -low matrix is 0. ■

Write $L = I + N$ and $L' = I + N'$. Since L and L' are unit lower triangular, N and N' are 1-low.

For (a), $LL' = I + N + N' + NN'$ by Theorem 0.9.4. Here $N + N'$ is 1-low, since a sum of entries that are both 0 is 0, and NN' is 2-low, hence 1-low, by the Claim. So $LL' - I$ is 1-low, which says that LL' is unit lower triangular.

For (b), let $K = I - N + N^2 - \dots + (-1)^{n-1}N^{n-1}$. Multiplying out, the sum telescopes:

$$(I + N)K = I + (-1)^{n-1}N^n = I,$$

where $N^n = 0$ because N^n is n -low by the Claim applied $n - 1$ times. The same computation gives $K(I + N) = I$, since N commutes with its own powers. Hence L is invertible with $L^{-1} = K$. Finally $K - I = \sum_{p=1}^{n-1} (-1)^p N^p$ is a sum of 1-low matrices, by the Claim, so K is unit lower triangular. This proves the lemma. ■

Transposing everything, the same statements hold for **unit upper triangular** matrices (upper triangular with 1's on the diagonal).

The next lemma decides invertibility of a triangular matrix from its diagonal, without any computation beyond back substitution.

Lemma 2.6.2 — Invertibility of a triangular matrix

An upper triangular matrix $U \in M_n(F)$ is invertible if and only if **every** diagonal entry $u_{11}, \dots, u_{nn}$ is non-zero.

Proof. ($\Leftarrow$) Suppose every $u_{kk} \neq 0$, and let $U\mathbf{x} = \mathbf{0}$. We show $x_k = 0$ for $k = n, n-1, \dots, 1$ in turn. Suppose $x_{k+1} = \dots = x_n = 0$ (nothing is assumed when $k = n$). Row k of $U\mathbf{x} = \mathbf{0}$ reads $u_{kk}x_k + \sum_{j>k} u_{kj}x_j = 0$, because $u_{kj} = 0$ for $j < k$. The sum vanishes, so $u_{kk}x_k = 0$, and $x_k = 0$ since $u_{kk} \neq 0$. Hence $\mathbf{x} = \mathbf{0}$ is the only solution of $U\mathbf{x} = \mathbf{0}$, and U is invertible by Theorem 2.4.4.

($\Rightarrow$) We prove the contrapositive. Suppose $u_{kk} = 0$ for some k . For each $j \leq k$, the entries of column j in rows $i \geq k$ are zero: if $i > j$ this is upper triangularity, and the only remaining case $i = j = k$ is $u_{kk} = 0$. So the first k columns of U lie in $W = \text{span}(\mathbf{e}_1, \dots, \mathbf{e}_{k-1})$. The list $(\mathbf{e}_1, \dots, \mathbf{e}_{k-1})$ is independent and spans W , so $\dim W = k-1$ (for $k=1$, $W = \{\mathbf{0}\}$ has dimension 0). By Theorem 1.6.3 (a), these k columns are linearly dependent, so the columns of U are dependent, and U is not invertible by Theorem 2.4.4. This proves the lemma. ■

2.6.4 Elimination without row swaps

To prove that elimination always produces an LU factorization, we first need to say exactly which elimination we mean. It is Gaussian elimination restricted to one kind of operation: subtracting a multiple of a row from a row **below** it.

Elimination without row swaps. Input: $A \in M_n(F)$. Put $A^{(0)} = A$. For $k = 1, 2, \dots, n-1$, write $A^{(k-1)} = (a_{ij})$ and do the following.

- If $a_{kk} \neq 0$, set $\ell_{ik} = a_{ik}/a_{kk}$ for $i = k+1, \dots, n$.
- If $a_{kk} = 0$ and $a_{ik} = 0$ for **every** $i > k$, set $\ell_{ik} = 0$ for $i = k+1, \dots, n$.
- If $a_{kk} = 0$ but $a_{ik} \neq 0$ for **some** $i > k$, **stop**: a row swap is needed.

If we did not stop, obtain $A^{(k)}$ from $A^{(k-1)}$ by replacing row i with $(\text{row } i) - \ell_{ik} \cdot (\text{row } k)$, for each $i = k+1, \dots, n$.

If the loop finishes without stopping, we say **elimination without row swaps succeeds** on A . The numbers ℓ_{ik} ($i > k$) are its **multipliers**, and $U = A^{(n-1)}$ is its **result**.

In both non-stopping cases the new (i, k) -entry is $a_{ik} - \ell_{ik}a_{kk} = 0$: in the first case by the choice of ℓ_{ik} , in the second because both terms are already 0. The multipliers are well defined: row k is not changed during step k , so it does not matter in which order the rows $i > k$ are updated.

The proof will write each step as a matrix. By Definition 2.4.1, the elementary matrix of the operation $R_i \rightarrow R_i + cR_k$ ($i \neq k$) is $I + cE_{ik}$, where E_{ik} is the matrix unit with 1 in entry (i, k) . We will write the matrix unit as a product of a column and a row,

$$E_{ik} = \mathbf{e}_i \mathbf{e}_k^T,$$

which holds because $\mathbf{e}_i \mathbf{e}_k^T$ has (p, q) -entry $(\mathbf{e}_i)_p (\mathbf{e}_k)_q$, equal to 1 if $p = i$ and $q = k$, and 0 otherwise. We will use one more rule constantly: for a vector $\mathbf{v} \in F^n$, the 1×1 matrix $\mathbf{e}_k^T \mathbf{v}$ is the k -th entry v_k , which we treat as a scalar. In particular $\mathbf{e}_k^T \mathbf{e}_j = 0$ for $j \neq k$.

For $1 \leq k \leq n$ and a vector $\mathbf{v} \in F^n$ whose first k entries are zero, put

$$G_k(\mathbf{v}) = I + \mathbf{v} \mathbf{e}_k^T.$$

Its only non-identity entries are v_i in position (i, k) , for $i > k$: it is the identity with the column vector $\mathbf{v}$ written into column k below the diagonal. So $G_k(\mathbf{v})$ is unit lower triangular.

Lemma 2.6.3 — Multiplier matrices

Let $1 \leq k \leq n$, and let $\mathbf{v}, \mathbf{v}_1, \dots, \mathbf{v}_m \in F^n$ be vectors such that the first k entries of $\mathbf{v}$ are zero and, for each j , the first j entries of $\mathbf{v}_j$ are zero ($m \leq n$). Then:

- (a) $G_k(\mathbf{v})$ is the product, in any order, of the elementary matrices $I + v_i E_{ik}$ for $i = k+1, \dots, n$;
- (b) $G_k(\mathbf{v})G_k(-\mathbf{v}) = G_k(-\mathbf{v})G_k(\mathbf{v}) = I$;
- (c) $G_1(\mathbf{v}_1)G_2(\mathbf{v}_2) \cdots G_m(\mathbf{v}_m) = I + \mathbf{v}_1 \mathbf{e}_1^\top + \mathbf{v}_2 \mathbf{e}_2^\top + \cdots + \mathbf{v}_m \mathbf{e}_m^\top$.

Proof. For (a), let $i, j > k$. By Theorem 0.9.4,

$$(I + a \mathbf{e}_i \mathbf{e}_k^\top)(I + b \mathbf{e}_j \mathbf{e}_k^\top) = I + a \mathbf{e}_i \mathbf{e}_k^\top + b \mathbf{e}_j \mathbf{e}_k^\top + ab \mathbf{e}_i (\mathbf{e}_k^\top \mathbf{e}_j) \mathbf{e}_k^\top,$$

and the last term is zero because $\mathbf{e}_k^\top \mathbf{e}_j = 0$ for $j \neq k$. Repeating this, one factor at a time, the product of the $I + v_i E_{ik} = I + v_i \mathbf{e}_i \mathbf{e}_k^\top$ in any order is $I + \sum_{i>k} v_i \mathbf{e}_i \mathbf{e}_k^\top = I + \mathbf{v} \mathbf{e}_k^\top$, since $\mathbf{v} = \sum_{i>k} v_i \mathbf{e}_i$.

For (b), $(I + \mathbf{v} \mathbf{e}_k^\top)(I - \mathbf{v} \mathbf{e}_k^\top) = I - \mathbf{v} (\mathbf{e}_k^\top \mathbf{v}) \mathbf{e}_k^\top = I$, because $\mathbf{e}_k^\top \mathbf{v} = v_k = 0$. Swapping the signs gives the other order.

For (c), we use induction on m ; the case $m = 1$ is the definition. Suppose the formula holds for $m - 1$. Then

$$\left(I + \sum_{j<m} \mathbf{v}_j \mathbf{e}_j^\top\right)(I + \mathbf{v}_m \mathbf{e}_m^\top) = I + \sum_{j<m} \mathbf{v}_j \mathbf{e}_j^\top + \mathbf{v}_m \mathbf{e}_m^\top + \sum_{j<m} \mathbf{v}_j (\mathbf{e}_j^\top \mathbf{v}_m) \mathbf{e}_m^\top.$$

For $j < m$, $\mathbf{e}_j^\top \mathbf{v}_m$ is entry j of $\mathbf{v}_m$, which is zero because the first m entries of $\mathbf{v}_m$ are zero. So the last sum vanishes, and the formula holds for m . This proves the lemma. ■

Part (c) is the heart of the matter: when the factors come in the order $G_1, G_2, \dots$, they **do not interfere**, and each vector $\mathbf{v}_j$ lands untouched in column j .

Theorem 2.6.4 — LU from elimination

Let $A \in M_n(F)$, and suppose elimination without row swaps succeeds on A , with multipliers ℓ_{ik} and result U . Let L be the unit lower triangular matrix whose (i, k) -entry is ℓ_{ik} for every $i > k$. Then U is upper triangular and

$$A = LU.$$

In particular A has an LU factorization.

Idea. Each step of elimination is left multiplication by a product of elementary matrices, which collapses to one matrix $G_k(-\ell_k)$. So $U = G_{n-1}(-\ell_{n-1}) \cdots G_1(-\ell_1)A$. To get A back, undo the steps in reverse order. Their inverses $G_k(\ell_k)$ then appear in the order

$G_1, G_2, \dots$, which is exactly the order in which, by [Lemma 2.6.3](#) (c), the multipliers drop into place.

Proof. Let $A^{(0)}, \dots, A^{(n-1)} = U$ be the matrices produced by the algorithm. For each k , let $\ell_k = \sum_{i>k} \ell_{ik} \mathbf{e}_i \in F^n$, a vector whose first k entries are zero.

Step 1: U is upper triangular. We show by induction on k that in $A^{(k)}$ every entry below the diagonal in columns $1, \dots, k$ is zero. For $k = 0$ there is nothing to prove. Suppose it holds for $A^{(k-1)}$. Step k changes only rows $i > k$, replacing row i by row i minus a multiple of row k . In columns $j < k$, the entry of row k lies below the diagonal ($k > j$), so it is 0 by the induction hypothesis, and the entries of row i in these columns stay 0. In column k , the new entries below the diagonal are 0, as noted after the algorithm. This completes the induction, and for $k = n - 1$ it says that U is upper triangular.

Step 2: each step is a matrix. Step k performs the operations “add $-\ell_{ik}$ times row k to row i ” for $i = k + 1, \dots, n$. By Theorem 2.4.1, applying them one after another multiplies $A^{(k-1)}$ on the left by the corresponding elementary matrices $I - \ell_{ik} E_{ik}$. By [Lemma 2.6.3](#) (a), their product is $G_k(-\ell_k)$. Hence $A^{(k)} = G_k(-\ell_k) A^{(k-1)}$, and therefore

$$U = G_{n-1}(-\ell_{n-1}) \cdots G_2(-\ell_2) G_1(-\ell_1) A.$$

Step 3: undo the steps. Multiply both sides on the left by $G_{n-1}(\ell_{n-1})$, then by $G_{n-2}(\ell_{n-2})$, and so on down to $G_1(\ell_1)$. By [Lemma 2.6.3](#) (b), each one cancels the factor next to it, and we get

$$G_1(\ell_1) G_2(\ell_2) \cdots G_{n-1}(\ell_{n-1}) U = A.$$

By [Lemma 2.6.3](#) (c), the product on the left of U is $I + \sum_{k=1}^{n-1} \ell_k \mathbf{e}_k^T$. Its (i, k) -entry for $i > k$ is ℓ_{ik} ; its diagonal entries are 1; its entries above the diagonal are 0. That is, it is L . Hence $A = LU$ with L unit lower triangular and U upper triangular, as claimed. ■

So the multipliers are not something to be remembered separately: they **are** the lower triangular factor. When a computer performs elimination it stores each ℓ_{ik} in the slot (i, k) that it has just cleared, and at the end the one array holds L below the diagonal and U on and above it.

Warning. The matrix that performs elimination, L^{-1} , is not L with its multipliers negated. In [Example \(An LU factorization from elimination\)](#) the operations multiply A on the left by

$$G_2(-\ell_2) G_1(-\ell_1) = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & -3 & 1 \end{pmatrix} \begin{pmatrix} 1 & 0 & 0 \\ -2 & 1 & 0 \\ 1 & 0 & 1 \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ -2 & 1 & 0 \\ 7 & -3 & 1 \end{pmatrix},$$

and the entry $7 = 1 + (-3)(-2)$ mixes two multipliers. This matrix is L^{-1} , by Step 3 of the proof. It is L , whose factors come in the order G_1, G_2 , that has the multipliers cleanly in place.

Quick check. Find the LU factorization of $A = \begin{pmatrix} 2 & 4 \\ 6 & 5 \end{pmatrix}$ over $\mathbb{Q}$ by elimination, and check

it by multiplying.

Solution. The pivot is 2 and the multiplier is $\ell_{21} = 6/2 = 3$. Subtracting 3 times row 1 from row 2 gives $U = \begin{pmatrix} 2 & 4 \\ 0 & 5-12 \end{pmatrix} = \begin{pmatrix} 2 & 4 \\ 0 & -7 \end{pmatrix}$, and $L = \begin{pmatrix} 1 & 0 \\ 3 & 1 \end{pmatrix}$. Check:
 $LU = \begin{pmatrix} 2 & 4 \\ 6 & 12-7 \end{pmatrix} = A$.

2.6.5 Solving with an LU factorization

Once $A = LU$ is known, each system $A\mathbf{x} = \mathbf{b}$ costs two substitutions.

Example. (Two right-hand sides, one factorization)

Let $A = \begin{pmatrix} 2 & 1 & -1 \\ 4 & 5 & -1 \\ -2 & 8 & 2 \end{pmatrix}$ with the factorization $A = LU$ of [Example \(An LU factorization from elimination\)](#). Solve $A\mathbf{x} = \mathbf{b}$ for

$$\mathbf{b} = \begin{pmatrix} -1 \\ -3 \\ -6 \end{pmatrix} \quad \text{and} \quad \mathbf{b}' = \begin{pmatrix} 3 \\ 7 \\ -2 \end{pmatrix}.$$

Solution. First right-hand side. Forward substitution in $L\mathbf{y} = \mathbf{b}$:

$$y_1 = -1, \quad 2y_1 + y_2 = -3 \Rightarrow y_2 = -1, \quad -y_1 + 3y_2 + y_3 = -6 \Rightarrow y_3 = -6 - 1 + 3 = -4.$$

Back substitution in $U\mathbf{x} = \mathbf{y}$:

$$-2x_3 = -4 \Rightarrow x_3 = 2, \quad 3x_2 + x_3 = -1 \Rightarrow x_2 = -1, \quad 2x_1 + x_2 - x_3 = -1 \Rightarrow x_1 = 1.$$

So $\mathbf{x} = (1, -1, 2)$. Check: $A\mathbf{x} = (2 - 1 - 2, 4 - 5 - 2, -2 - 8 + 4) = (-1, -3, -6) = \mathbf{b}$.

Second right-hand side. $L\mathbf{y} = \mathbf{b}'$: $y_1 = 3, y_2 = 7 - 2 \cdot 3 = 1, y_3 = -2 + y_1 - 3y_2 = -2 + 3 - 3 = -2$. Then $U\mathbf{x} = \mathbf{y}$: $-2x_3 = -2$, so $x_3 = 1$; $3x_2 + 1 = 1$, so $x_2 = 0$; $2x_1 + 0 - 1 = 3$, so $x_1 = 2$. So $\mathbf{x}' = (2, 0, 1)$, and indeed $A\mathbf{x}' = (4 + 0 - 1, 8 + 0 - 1, -4 + 0 + 2) = (3, 7, -2)$.

Since the diagonal of U is $2, 3, -2$, all non-zero, U is invertible by [Lemma 2.6.2](#), L is invertible by [Lemma 2.6.1](#), and so $A = LU$ is invertible. Each solution found is therefore the only one.

How much is saved? Count the arithmetic operations $(+, -, \times, \div)$ for an $n \times n$ matrix. At step k of elimination there are $n - k$ rows below the pivot. Each needs one division for its multiplier, then $n - k$ multiplications and $n - k$ subtractions in the columns to the right of column k . With $j = n - k$, the total is

$$\sum_{j=1}^{n-1} j(2j+1) = \frac{n(n-1)(4n+1)}{6} \approx \frac{2}{3}n^3.$$

Forward substitution with a unit lower triangular L uses $n(n-1)$ operations, and back substitution n^2 , so a pair of triangular solves costs $2n^2 - n$. For $n = 1000$ that is about 6.7×10^8

operations for the factorization, done once, against about 2×10^6 for each new right-hand side.

2.6.6 Uniqueness

The factorization of [Example \(An LU factorization from elimination\)](#) came from one particular procedure. Could another procedure give a different pair (L, U) ? For invertible matrices, no.

Theorem 2.6.5 — Uniqueness of the LU factorization

Let $A \in M_n(F)$ be **invertible**. If $A = L_1U_1$ and $A = L_2U_2$ are LU factorizations, then $L_1 = L_2$ and $U_1 = U_2$.

Idea. Take the difference in multiplicative form: from $L_1U_1 = L_2U_2$ we get $(L_2^{-1}L_1)U_1 = U_2$. The matrix $M = L_2^{-1}L_1$ is unit lower triangular, and multiplying U_1 by it must leave an upper triangular matrix. Clearing below the diagonal column by column, the non-zero pivots of U_1 force M to have nothing below its diagonal. Invertibility of A is what provides those non-zero pivots.

Proof. By [Lemma 2.6.1](#) (b), L_1 and L_2 are invertible. Hence $U_1 = L_1^{-1}A$ is a product of invertible matrices, so it is invertible by Theorem 0.9.7, and by [Lemma 2.6.2](#) its diagonal entries $u_{11}, \dots, u_{nn}$ are all non-zero.

Let $M = L_2^{-1}L_1$. By [Lemma 2.6.1](#), $M = (m_{ij})$ is unit lower triangular. Multiplying $L_1U_1 = L_2U_2$ on the left by L_2^{-1} gives $MU_1 = U_2$, so MU_1 is upper triangular. We show $m_{ij} = 0$ for all $i > j$, by induction on the column j .

Fix j , and suppose $m_{ik} = 0$ whenever $k < j$ and $i > k$. Let $i > j$. Since MU_1 is upper triangular and $i > j$,

$$0 = (MU_1)_{ij} = \sum_{k=1}^n m_{ik}u_{kj} = \sum_{k \leq j} m_{ik}u_{kj} = m_{ij}u_{jj},$$

where the third equality uses $u_{kj} = 0$ for $k > j$, and the last uses $m_{ik} = 0$ for $k < j$ (here $i > j > k$). Since $u_{jj} \neq 0$, we get $m_{ij} = 0$. This completes the induction, so M is lower triangular with 1's on the diagonal and nothing below it: $M = I$.

Hence $L_2^{-1}L_1 = I$, and multiplying by L_2 gives $L_1 = L_2$. Then $U_1 = L_1^{-1}A = L_2^{-1}A = U_2$. This proves the theorem. ■

Two consequences are worth saying. The LU factorization of an invertible matrix does not depend on how it was found, so it is legitimate to speak of **the** factors L and U . And if elimination without row swaps succeeds on an invertible A , then every pivot is non-zero: row k is not changed after step k , so the pivot of step k is u_{kk} , which is non-zero by the first paragraph of the proof. So the second case of the algorithm (a zero pivot over a zero column) never occurs.

Remark. The hypothesis “invertible” cannot be dropped. For every $c \in F$,

$$\begin{pmatrix} 0 & 0 \\ 0 & 1 \end{pmatrix} = \begin{pmatrix} 1 & 0 \\ c & 1 \end{pmatrix} \begin{pmatrix} 0 & 0 \\ 0 & 1 \end{pmatrix},$$

since the first column of the right factor is zero. Taking $c = 0$ and $c = 1$ gives two different LU factorizations. The proof breaks exactly at “ $u_{jj} \neq 0$ ”, here for $j = 1$.

2.6.7 Row swaps and the PLU factorization

Example (An invertible matrix with no LU factorization) showed that some matrices, even invertible ones, have no LU factorization, because elimination meets a zero pivot with a non-zero entry below it. The cure in Gaussian elimination is to swap rows. We now show that the swaps can all be collected in front, as a single matrix that rearranges the rows of A **before** elimination starts. After that rearrangement, no swaps are needed.

Matrices that rearrange rows have a name.

Definition 2.6.2 — Permutation matrix

Let $\sigma \in S_n$ be a permutation of $\{1, \dots, n\}$. The **permutation matrix** of σ is

$$P_\sigma = (\mathbf{e}_{\sigma(1)} \quad \mathbf{e}_{\sigma(2)} \quad \cdots \quad \mathbf{e}_{\sigma(n)}) \in M_n(F),$$

the matrix whose j -th column is $\mathbf{e}_{\sigma(j)}$. A matrix is a **permutation matrix** if it equals P_σ for some $\sigma \in S_n$.

Equivalently, a permutation matrix has exactly one 1 in each row and in each column, and zeros elsewhere. The identity is P_{id} . The matrix $\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$ of **Example (An invertible matrix with no LU factorization)** is P_τ for the transposition $\tau = (1 \ 2)$. By contrast $\begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}$ is **not** a permutation matrix: each column is still non-zero, but its first row contains two 1's, and its second column $(1, 1)$ is not a standard basis vector.

Lemma 2.6.6 — Permutation matrices

Let $\sigma, \tau \in S_n$.

- (a) $P_\sigma P_\tau = P_{\sigma \circ \tau}$. In particular a product of permutation matrices is a permutation matrix.
- (b) P_σ is invertible and $P_\sigma^{-1} = P_\sigma^\top$.
- (c) If $\tau = (r \ s)$ is a transposition, then P_τ is the elementary matrix P_{rs} that swaps rows r and s , and $P_\tau^2 = I$.
- (d) Let $1 \leq k < r < s \leq n$, let $\tau = (r \ s)$, and let $\mathbf{v} \in F^n$ have its first k entries zero. Then $P_\tau \mathbf{v}$ has its first k entries zero, and

$$P_\tau G_k(\mathbf{v}) P_\tau = G_k(P_\tau \mathbf{v}).$$

Proof.

- (a) By Theorem 0.9.3, $P_\tau \mathbf{e}_j = \mathbf{e}_{\tau(j)}$ and so $P_\sigma P_\tau \mathbf{e}_j = P_\sigma \mathbf{e}_{\tau(j)} = \mathbf{e}_{\sigma(\tau(j))}$. So the j -th column of $P_\sigma P_\tau$ is $\mathbf{e}_{(\sigma \circ \tau)(j)}$, for every j , which is the claim.
- (b) By Theorem 0.9.2, $(P_\sigma^\top P_\sigma)_{ij} = \mathbf{e}_{\sigma(i)}^\top \mathbf{e}_{\sigma(j)}$, which is 1 if $\sigma(i) = \sigma(j)$ and 0 otherwise. Since σ is injective, $\sigma(i) = \sigma(j)$ exactly when $i = j$. Hence $P_\sigma^\top P_\sigma = I$, and by Theorem 2.4.5, P_σ is invertible with inverse $P_\sigma^\top$.
- (c) Swapping rows r and s of I produces a matrix whose column r is $\mathbf{e}_s$, whose column s is $\mathbf{e}_r$, and whose other columns j are $\mathbf{e}_j$. These are the columns $\mathbf{e}_{\tau(j)}$, so the swap matrix is P_τ . Since $\tau \circ \tau = \text{id}$, part (a) gives $P_\tau^2 = P_{\text{id}} = I$.
- (d) Write $P = P_\tau$. Its effect on a vector is to exchange entries r and s , which are both beyond position k , so $P\mathbf{v}$ still has its first k entries zero. Since $k \neq r, s$, we have $P\mathbf{e}_k = \mathbf{e}_{\tau(k)} = \mathbf{e}_k$. By (b) and (c), $P^\top = P^{-1} = P$, so $\mathbf{e}_k^\top P = \mathbf{e}_k^\top P^\top = (P\mathbf{e}_k)^\top = \mathbf{e}_k^\top$ by Theorem 0.9.5. Therefore

$$P(I + \mathbf{v} \mathbf{e}_k^\top)P = P^2 + (P\mathbf{v})(\mathbf{e}_k^\top P) = I + (P\mathbf{v}) \mathbf{e}_k^\top = G_k(P\mathbf{v}),$$

using $P^2 = I$ from (c). This proves the lemma. ■

By (c) and Theorem 2.4.1, $P_\tau A$ is A with rows r and s swapped, and by (a) every permutation matrix is a product of such swaps applied in turn, since for $n \geq 2$ every permutation is a product of transpositions (Theorem 0.10.4). So PA is always A with its rows rearranged.

Definition 2.6.3 — PLU factorization

Let $A \in M_n(F)$. A **PLU factorization** of A is a triple (P, L, U) of matrices in $M_n(F)$ such that

$$PA = LU, \quad P \text{ is a permutation matrix,} \quad L \text{ is unit lower triangular,} \quad U \text{ is upper triangular.}$$

In words: after rearranging the rows of A by P , the matrix has an LU factorization. By Lemma 2.6.6 (b) this can also be written $A = P^\top LU$, which is why some books write the factorization as $A = PLU$ with their P our $P^\top$. For example, $A = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$ has the PLU factorization (A, I, I) , since $AA = I = I \cdot I$. An LU factorization is the special case $P = I$.

To produce P , run elimination and allow a swap whenever a zero pivot has a non-zero entry below it.

Elimination with row swaps. Input: $A \in M_n(F)$. Put $A^{(0)} = A$. For $k = 1, 2, \dots, n-1$, with $A^{(k-1)} = (a_{ij})$:

- If $a_{ik} = 0$ for **every** $i \geq k$, set $P_k = I$ and all $\ell_{ik} = 0$, and let $A^{(k)} = A^{(k-1)}$.
- Otherwise choose some $r_k \geq k$ with $a_{r_k k} \neq 0$. Let P_k be the matrix swapping rows k and r_k (so $P_k = I$ if $r_k = k$), and let $B = P_k A^{(k-1)}$. Set $\ell_{ik} = b_{ik}/b_{kk}$ for $i > k$, and obtain $A^{(k)}$ from B by replacing row i with (row i) $-\ell_{ik} \cdot$ (row k) for each $i > k$.

This version never stops. Here is the result it delivers.

Theorem 2.6.7 — Existence of a PLU factorization

Every $A \in M_n(F)$ has a PLU factorization.

Idea. Step roadmap. ① As before, each step is a matrix: $A^{(k)} = G_k(-\ell_k)P_k A^{(k-1)}$, and the final matrix U is upper triangular. ② The swaps are tangled with the multiplier matrices. A swap of two rows below row k can be pushed to the right past G_k at the price of swapping the entries of the multiplier vector: this is Lemma 2.6.6 (d). ③ Once all swaps sit next to A , the rest is the LU proof word for word. In practice, ② says: when you swap two rows, also swap the multipliers already stored in those rows.

Proof. Run elimination with row swaps on A , producing $A^{(0)}, \dots, A^{(n-1)} = U$, swap matrices $P_1, \dots, P_{n-1}$ and multipliers ℓ_{ik} . Let $\ell_k = \sum_{i>k} \ell_{ik} \mathbf{e}_i$, whose first k entries are zero.

Step 1. Exactly as in Step 1 of the proof of Theorem 2.6.4, every entry of $A^{(k)}$ below the diagonal in columns $1, \dots, k$ is zero. The only new point is the swap in step k : it exchanges rows k and $r_k \geq k$, whose entries in columns $j < k$ are below the diagonal and hence zero, so the swap does not disturb these columns. After the swap the pivot b_{kk} is non-zero, and the new entries below it are $b_{ik} - \ell_{ik}b_{kk} = 0$. Hence U is upper triangular. By Lemma 2.6.6 (c), Theorem 2.4.1 and Lemma 2.6.3 (a), as in Step 2 of that proof,

$$A^{(k)} = G_k(-\ell_k) P_k A^{(k-1)} \quad (k = 1, \dots, n-1).$$

This also holds in the first case of the algorithm, where $P_k = I$ and $G_k(\mathbf{0}) = I$.

Step 2. For $1 \leq k \leq m \leq n-1$ put $\ell_k^{(m)} = P_m P_{m-1} \cdots P_{k+1} \ell_k$, with $\ell_m^{(m)} = \ell_m$. We claim that the first k entries of $\ell_k^{(m)}$ are zero and that

$$G_m(-\ell_m) P_m \cdots G_1(-\ell_1) P_1 = G_m(-\ell_m^{(m)}) \cdots G_1(-\ell_1^{(m)}) P_m \cdots P_1.$$

For $m = 1$ both sides are $G_1(-\ell_1) P_1$. Suppose the claim holds for m , and multiply it on the left by $G_{m+1}(-\ell_{m+1}) P_{m+1}$. The matrix P_{m+1} is I or swaps rows $m+1$ and $r_{m+1} > m+1$, both beyond every $k \leq m$. So by Lemma 2.6.6 (c) and (d), for each $k \leq m$ and each vector $\mathbf{v}$ with its first k entries zero,

$$P_{m+1} G_k(\mathbf{v}) = P_{m+1} G_k(\mathbf{v}) P_{m+1} P_{m+1} = G_k(P_{m+1} \mathbf{v}) P_{m+1},$$

and $P_{m+1} \mathbf{v}$ again has its first k entries zero. Applying this for $k = m, m-1, \dots, 1$ moves P_{m+1} to the right past every factor G_k , replacing $\ell_k^{(m)}$ by $P_{m+1} \ell_k^{(m)} = \ell_k^{(m+1)}$ (note $P_{m+1}(-\mathbf{v}) = -P_{m+1} \mathbf{v}$). This is the claim for $m+1$.

Step 3. By Step 1, $U = G_{n-1}(-\ell_{n-1}) P_{n-1} \cdots G_1(-\ell_1) P_1 A$. Write $\ell'_k = \ell_k^{(n-1)}$ and $P = P_{n-1} \cdots P_1$. By Step 2,

$$U = G_{n-1}(-\ell'_{n-1}) \cdots G_1(-\ell'_1) P A.$$

P is a permutation matrix by Lemma 2.6.6 (a) and (c). Multiplying on the left by $G_{n-1}(\ell'_{n-1})$, then $G_{n-2}(\ell'_{n-2})$, and so on, and using Lemma 2.6.3 (b) and (c) as in Step

3 of the proof of [Theorem 2.6.4](#),

$$PA = LU, \quad L = I + \sum_{k=1}^{n-1} \ell'_k \mathbf{e}_k^T,$$

where L is unit lower triangular because the first k entries of each ℓ'_k are zero. So (P, L, U) is a PLU factorization of A . This proves the theorem. ■

Example. (A swap in the middle)

Find a PLU factorization of

$$A = \begin{pmatrix} 1 & 2 & 1 \\ 2 & 4 & 3 \\ 3 & 7 & 2 \end{pmatrix} \in M_3(\mathbb{Q}).$$

Solution. Step 1. The pivot 1 is non-zero, so no swap. The multipliers are $\ell_{21} = 2$ and $\ell_{31} = 3$:

$$A \xrightarrow[R_3 \rightarrow R_3 - 3R_1]{R_2 \rightarrow R_2 - 2R_1} \begin{pmatrix} 1 & 2 & 1 \\ 0 & 0 & 1 \\ 0 & 1 & -1 \end{pmatrix}, \quad \ell_1 = \begin{pmatrix} 0 \\ 2 \\ 3 \end{pmatrix}.$$

Step 2. The pivot position $(2, 2)$ holds 0, but the entry below it is 1. Elimination without row swaps would stop here. We swap rows 2 and 3, so $P_2 = P_{23}$:

$$\xrightarrow{R_2 \leftrightarrow R_3} \begin{pmatrix} 1 & 2 & 1 \\ 0 & 1 & -1 \\ 0 & 0 & 1 \end{pmatrix} = U.$$

The entry below the new pivot is already 0, so $\ell_{32} = 0$, and U is upper triangular.

By Step 2 of the proof, the stored multipliers of step 1 are swapped too: $\ell'_1 = P_2 \ell_1 = (0, 3, 2)$. Hence

$$P = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{pmatrix}, \quad L = \begin{pmatrix} 1 & 0 & 0 \\ 3 & 1 & 0 \\ 2 & 0 & 1 \end{pmatrix}.$$

Check: PA is A with rows 2 and 3 swapped, and

$$LU = \begin{pmatrix} 1 & 2 & 1 \\ 3 & 6+1 & 3-1 \\ 2 & 4 & 2+1 \end{pmatrix} = \begin{pmatrix} 1 & 2 & 1 \\ 3 & 7 & 2 \\ 2 & 4 & 3 \end{pmatrix} = PA.$$

Notice where the multipliers ended up: 3 in row 2 and 2 in row 3, traveling with the rows they belong to. To solve $A\mathbf{x} = \mathbf{b}$, multiply by P and solve $LU\mathbf{x} = P\mathbf{b}$ by two substitutions; since P is invertible, this changes no solutions. For $\mathbf{b} = (2, 3, 8)$, $P\mathbf{b} = (2, 8, 3)$. Forward substitution in $L\mathbf{y} = P\mathbf{b}$ gives $y_1 = 2$, $y_2 = 8 - 3 \cdot 2 = 2$, $y_3 = 3 - 2 \cdot 2 = -1$. Back substitution in $U\mathbf{x} = \mathbf{y}$ gives $x_3 = -1$, $x_2 = 2 + x_3 = 1$, $x_1 = 2 - 2x_2 - x_3 = 1$. Check: $A(1, 1, -1) = (1 + 2 - 1, 2 + 4 - 3, 3 + 7 - 2) = (2, 3, 8)$.

Remark. Over $\mathbb{R}$, computers do not choose just any non-zero pivot. They choose the entry of **largest absolute value** in the column, which is called **partial pivoting**. The reason is rounding. For $A = \begin{pmatrix} \varepsilon & 1 \\ 1 & 1 \end{pmatrix}$ with $\varepsilon = 10^{-20}$, elimination without swaps gives the multiplier 10^{20} and $u_{22} = 1 - 10^{20}$, which in standard double-precision arithmetic is stored as -10^{20} . Multiplying back, the computed LU has $(2, 2)$ -entry $10^{20} - 10^{20} = 0$ instead of 1. Swapping the rows first gives the harmless multiplier ε . Chapter 23 makes this precise.

2.6.8 Symmetric matrices: the LDL factorization

When A is symmetric, the two factors of $A = LU$ are secretly related: up to a diagonal matrix, U is the transpose of L . In [Example \(An LU factorization from elimination\)](#) the factors look unrelated, but that matrix is not symmetric. Try a symmetric one:

$$\begin{pmatrix} 1 & 3 & -2 \\ 3 & 7 & -8 \\ -2 & -8 & 5 \end{pmatrix} \xrightarrow[R_3 \rightarrow R_3 + 2R_1]{R_2 \rightarrow R_2 - 3R_1} \begin{pmatrix} 1 & 3 & -2 \\ 0 & -2 & -2 \\ 0 & -2 & 1 \end{pmatrix} \xrightarrow{R_3 \rightarrow R_3 - R_2} \begin{pmatrix} 1 & 3 & -2 \\ 0 & -2 & -2 \\ 0 & 0 & 3 \end{pmatrix} = U,$$

with multipliers $\ell_{21} = 3$, $\ell_{31} = -2$, $\ell_{32} = 1$. Divide each row of U by its diagonal entry: the rows become $(1, 3, -2)$, $(0, 1, 1)$, $(0, 0, 1)$. These are the columns of $L = \begin{pmatrix} 1 & 0 & 0 \\ 3 & 1 & 0 \\ -2 & 1 & 1 \end{pmatrix}$, read as rows. So $U = DL^T$ with $D = \text{diag}(1, -2, 3)$, and $A = LDL^T$.

Theorem 2.6.8 — LDL factorization

Let $A \in M_n(F)$ be **symmetric** and **invertible**, and suppose $A = LU$ is an LU factorization. Then the diagonal entries $u_{11}, \dots, u_{nn}$ of U are non-zero, and with $D = \text{diag}(u_{11}, \dots, u_{nn})$,

$$U = DL^T \quad \text{and} \quad A = LDL^T.$$

Idea. Pull the diagonal out of U , writing $U = DV$ with V unit upper triangular, so $A = LDV$. Transposing, $A = A^T = V^T(DL^T)$, which is a second LU factorization of the same invertible matrix. Uniqueness forces the two to agree, so $L = V^T$.

Proof. As in the proof of [Theorem 2.6.5](#), $U = L^{-1}A$ is invertible, so its diagonal entries are non-zero by [Lemma 2.6.2](#). Hence D is invertible, with $D^{-1} = \text{diag}(u_{11}^{-1}, \dots, u_{nn}^{-1})$. Let $V = D^{-1}U$. Row i of V is u_{ii}^{-1} times row i of U , so V is upper triangular with $v_{ii} = u_{ii}^{-1}u_{ii} = 1$, and $A = LDV$.

Since A is symmetric and $D^T = D$, Theorem 0.9.5 gives

$$A = A^T = (LDV)^T = V^T DL^T = V^T (DL^T).$$

Here V^T is unit lower triangular, because V is unit upper triangular. And DL^T is upper triangular: its (i, j) -entry is $u_{ij}\ell_{ji}$, and $\ell_{ji} = 0$ when $j < i$, since L is lower triangular. So (V^T, DL^T) is an LU factorization of the invertible matrix A . By [Theorem 2.6.5](#), $V^T = L$ and $DL^T = U$. Therefore $A = LU = LDL^T$, as claimed.

The factorization stores about half as much: L and the diagonal D determine everything. Over $\mathbb{R}$, when every entry of D is positive, one can split $D = D^{1/2}D^{1/2}$ and write $A = (LD^{1/2})(LD^{1/2})^T$; this is the **Cholesky factorization**, which Chapter 12 connects with positive definite matrices.

Remark. Symmetry alone does not give an LU factorization: $\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$ is symmetric and invertible, and has none by [Example \(An invertible matrix with no LU factorization\)](#). The theorem assumes the factorization exists and then improves it.

2.6.9 Exercises

2.6.9.1 A. Check your understanding

Exercise A1.

- Define an LU factorization of a matrix $A \in M_n(F)$, and a PLU factorization.
- True or false: every invertible matrix has an LU factorization. Justify your answer.
- Suppose $A = LU$. In which order are the two triangular systems solved to find a solution of $A\mathbf{x} = \mathbf{b}$, and what is each one?
- True or false: the product of two unit lower triangular $n \times n$ matrices is unit lower triangular. Justify your answer.
- True or false: the zero matrix in $M_2(F)$ has more than one LU factorization. Justify your answer.
- What is the inverse of a permutation matrix P ?

Solution.

- See [Definition 2.6.1](#): $A = LU$ with L unit lower triangular and U upper triangular. See [Definition 2.6.3](#): $PA = LU$ with P a permutation matrix, L unit lower triangular and U upper triangular.
- False. $\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$ is invertible but has no LU factorization, by [Example \(An invertible matrix with no LU factorization\)](#).
- First $L\mathbf{y} = \mathbf{b}$, by forward substitution; then $U\mathbf{x} = \mathbf{y}$, by back substitution. Then $A\mathbf{x} = L(U\mathbf{x}) = L\mathbf{y} = \mathbf{b}$.
- True, by [Lemma 2.6.1](#) (a).
- True. For every $c \in F$, $\begin{pmatrix} 1 & 0 \\ c & 1 \end{pmatrix} \begin{pmatrix} 0 \\ 0 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \end{pmatrix}$, so $\left(\begin{pmatrix} 1 & 0 \\ c & 1 \end{pmatrix}, \begin{pmatrix} 0 \\ 0 \end{pmatrix}\right)$ is an LU factorization; $c = 0$ and $c = 1$ give two different ones. This does not contradict [Theorem 2.6.5](#), because 0 is not invertible.
- $P^{-1} = P^T$, by [Lemma 2.6.6](#) (b).

2.6.9.2 B. Practice

Exercise B1: An LU factorization.

Find the LU factorization of

$$A = \begin{pmatrix} 1 & 2 & 0 \\ -3 & -7 & 4 \\ 2 & 6 & -3 \end{pmatrix} \in M_3(\mathbb{Q})$$

by elimination without row swaps. Hence show that A is invertible.

Solution. Step 1. The pivot is 1, with multipliers $\ell_{21} = -3$ and $\ell_{31} = 2$:

$$A \xrightarrow[R_3 \rightarrow R_3 - 2R_1]{R_2 \rightarrow R_2 - (-3)R_1} \begin{pmatrix} 1 & 2 & 0 \\ 0 & -1 & 4 \\ 0 & 2 & -3 \end{pmatrix}.$$

Step 2. The pivot is -1 , with multiplier $\ell_{32} = 2/(-1) = -2$:

$$\xrightarrow{R_3 \rightarrow R_3 - (-2)R_2} \begin{pmatrix} 1 & 2 & 0 \\ 0 & -1 & 4 \\ 0 & 0 & 5 \end{pmatrix} = U.$$

By [Theorem 2.6.4](#),

$$L = \begin{pmatrix} 1 & 0 & 0 \\ -3 & 1 & 0 \\ 2 & -2 & 1 \end{pmatrix}, \quad A = LU.$$

$$\text{Check: } LU = \begin{pmatrix} 1 & 2 & 0 \\ -3 & -6 - 1 & 4 \\ 2 & 4 + 2 & -8 + 5 \end{pmatrix} = A.$$

The diagonal entries 1, -1 , 5 of U are non-zero, so U is invertible by [Lemma 2.6.2](#), and L is invertible by [Lemma 2.6.1](#). Hence $A = LU$ is invertible by Theorem 0.9.7.

Exercise B2: Two systems, one factorization.

With $A = LU$ as in [Exercise \(B1: An LU factorization\)](#), solve $A\mathbf{x} = \mathbf{b}$ for $\mathbf{b} = (3, -6, 5)$ and for $\mathbf{b} = (0, 1, -2)$, using only forward and back substitution.

Solution. For $\mathbf{b} = (3, -6, 5)$. $L\mathbf{y} = \mathbf{b}$: $y_1 = 3$; $-3y_1 + y_2 = -6$, so $y_2 = 3$; $2y_1 - 2y_2 + y_3 = 5$, so $y_3 = 5 - 6 + 6 = 5$. Then $U\mathbf{x} = \mathbf{y}$: $5x_3 = 5$, so $x_3 = 1$; $-x_2 + 4x_3 = 3$, so $x_2 = 1$; $x_1 + 2x_2 = 3$, so $x_1 = 1$. Hence $\mathbf{x} = (1, 1, 1)$. Check: $A\mathbf{x} = (1+2, -3-7+4, 2+6-3) = (3, -6, 5)$.

For $\mathbf{b} = (0, 1, -2)$. $L\mathbf{y} = \mathbf{b}$: $y_1 = 0$; $y_2 = 1 + 3y_1 = 1$; $y_3 = -2 - 2y_1 + 2y_2 = 0$. Then $U\mathbf{x} = \mathbf{y}$: $5x_3 = 0$, so $x_3 = 0$; $-x_2 + 0 = 1$, so $x_2 = -1$; $x_1 + 2(-1) = 0$, so $x_1 = 2$. Hence $\mathbf{x} = (2, -1, 0)$. Check: $A\mathbf{x} = (2-2, -6+7, 4-6) = (0, 1, -2)$.

Since A is invertible ([Exercise \(B1: An LU factorization\)](#)), each solution is the only one, by Theorem 0.9.7.

Exercise B3: A PLU factorization.

Let $A = \begin{pmatrix} 0 & 1 & 2 \\ 1 & 1 & 1 \\ 2 & 3 & 1 \end{pmatrix} \in M_3(\mathbb{Q})$.

- (a) Explain why A has no LU factorization.
- (b) Find a PLU factorization of A .
- (c) Hence solve $A\mathbf{x} = (0, 2, 7)$.

Solution.

- (a) If $A = LU$, then the $(1,1)$ -entry gives $u_{11} = a_{11} = 0$, and the $(2,1)$ -entry gives $a_{21} = \ell_{21}u_{11} = 0$. But $a_{21} = 1 \neq 0$, a contradiction, exactly as in [Example \(An invertible matrix with no LU factorization\)](#).
- (b) The first pivot position holds 0 and the entry below is 1, so swap rows 1 and 2, $P_1 = P_{12}$:

$$A \xrightarrow{R_1 \leftrightarrow R_2} \begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 2 \\ 2 & 3 & 1 \end{pmatrix} \xrightarrow[\substack{R_2 \rightarrow R_2 - 0R_1 \\ R_3 \rightarrow R_3 - 2R_1}]{R_2 \rightarrow R_2 - 0R_1} \begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 2 \\ 0 & 1 & -1 \end{pmatrix} \xrightarrow{R_3 \rightarrow R_3 - 1R_2} \begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 2 \\ 0 & 0 & -3 \end{pmatrix} = U.$$

The multipliers are $\ell_{21} = 0$, $\ell_{31} = 2$, $\ell_{32} = 1$, and no swap happens after step 1, so they need no rearranging. Hence

$$P = \begin{pmatrix} 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{pmatrix}, \quad L = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 2 & 1 & 1 \end{pmatrix}.$$

$$\text{Check: } PA = \begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 2 \\ 2 & 3 & 1 \end{pmatrix}, \text{ and } LU = \begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 2 \\ 2 & 2+1 & 2+2-3 \end{pmatrix} = PA.$$

- (c) $A\mathbf{x} = \mathbf{b}$ holds if and only if $PA\mathbf{x} = P\mathbf{b}$, since P is invertible ([Lemma 2.6.6 \(b\)](#)). Swapping the first two entries of $\mathbf{b} = (0, 2, 7)$, we solve $LU\mathbf{x} = P\mathbf{b} = (2, 0, 7)$. $Ly = (2, 0, 7)$: $y_1 = 2$, $y_2 = 0$, $y_3 = 7 - 2 \cdot 2 - 0 = 3$. $U\mathbf{x} = \mathbf{y}$: $-3x_3 = 3$, so $x_3 = -1$; $x_2 + 2x_3 = 0$, so $x_2 = 2$; $x_1 + x_2 + x_3 = 2$, so $x_1 = 1$. Hence $\mathbf{x} = (1, 2, -1)$. Check: $A\mathbf{x} = (0 + 2 - 2, 1 + 2 - 1, 2 + 6 - 1) = (0, 2, 7)$.

2.6.9.3 C. Going deeper**Exercise C1: Leading principal submatrices.**

For $A \in M_n(F)$ and $1 \leq k \leq n$, the **leading principal submatrix** $A_k \in M_k(F)$ is the top-left $k \times k$ corner of A , with entries a_{ij} for $i, j \leq k$.

- (a) Let L be lower triangular and U upper triangular in $M_n(F)$. Show that $(LU)_k = L_k U_k$ for every k .
- (b) Deduce that if A has an LU factorization $A = LU$ with U invertible, then every leading principal submatrix $A_1, \dots, A_n$ is invertible.
- (c) Prove the converse: if $A_1, \dots, A_n$ are all invertible, then A has an LU factorization

with U invertible.

Hint for (c): induction on n . Write A with A_{n-1} in its top-left corner, a column $\mathbf{b}$, a row $\mathbf{c}^\top$ and a corner entry d , and look for L and U of the same shape.

Solution.

- (a) Let $i, j \leq k$. By Definition 0.9.9, $(LU)_{ij} = \sum_{l=1}^n \ell_{il} u_{lj}$. If $l > k$, then $l > i$, so $\ell_{il} = 0$ because L is lower triangular. Hence $(LU)_{ij} = \sum_{l=1}^k \ell_{il} u_{lj} = (L_k U_k)_{ij}$, as claimed.
- (b) Suppose $A = LU$ with U invertible. By (a), $A_k = L_k U_k$. The matrix L_k is unit lower triangular, so it is invertible by Lemma 2.6.1. By Lemma 2.6.2, $u_{11}, \dots, u_{nn}$ are non-zero; U_k is upper triangular with diagonal $u_{11}, \dots, u_{kk}$, so U_k is invertible by the same lemma. Hence A_k is invertible by Theorem 0.9.7.
- (c) We prove by induction on n that every $A \in M_n(F)$ whose leading principal submatrices are all invertible has an LU factorization with U invertible. For $n = 1$, $A = (a_{11})$ with $a_{11} \neq 0$, and $A = (1)(a_{11})$ works. Let $n \geq 2$ and suppose the statement holds for $n - 1$. Write

$$A = \begin{pmatrix} A_{n-1} & \mathbf{b} \\ \mathbf{c}^\top & d \end{pmatrix}, \quad \mathbf{b}, \mathbf{c} \in F^{n-1}, d \in F.$$

The leading principal submatrices of A_{n-1} are $A_1, \dots, A_{n-1}$, all invertible, so by the induction hypothesis $A_{n-1} = L'U'$ with L' unit lower triangular and U' upper triangular and invertible. Put $\mathbf{y} = L'^{-1}\mathbf{b}$, $\mathbf{x}^\top = \mathbf{c}^\top U'^{-1}$ and $z = d - \mathbf{x}^\top \mathbf{y}$, and let

$$L = \begin{pmatrix} L' & \mathbf{0} \\ \mathbf{x}^\top & 1 \end{pmatrix}, \quad U = \begin{pmatrix} U' & \mathbf{y} \\ \mathbf{0}^\top & z \end{pmatrix}.$$

Then L is unit lower triangular and U is upper triangular. Computing LU entry by entry, splitting each sum $\sum_{l=1}^n$ into $l \leq n-1$ and $l = n$: the top-left block is $L'U' + \mathbf{0}\mathbf{0}^\top = A_{n-1}$; the top-right column is $L'\mathbf{y} + \mathbf{0}z = \mathbf{b}$; the bottom-left row is $\mathbf{x}^\top U' + 1 \cdot \mathbf{0}^\top = \mathbf{c}^\top$; the corner is $\mathbf{x}^\top \mathbf{y} + z = d$. Hence $A = LU$.

Finally $A = A_n$ is invertible by hypothesis, and L is invertible by Lemma 2.6.1, so $U = L^{-1}A$ is invertible by Theorem 0.9.7. This completes the induction.

In the language of elimination: z is the n -th pivot, and (c) shows that invertible leading corners are exactly what keeps every pivot non-zero.

Exercise C2: An LDL factorization.

Let $A = \begin{pmatrix} 4 & -4 & 8 \\ -4 & 5 & -11 \\ 8 & -11 & 27 \end{pmatrix} \in M_3(\mathbb{R})$.

- (a) Find the LU factorization of A , show that A is invertible, and write $A = LDL^\top$.
 (b) Deduce that $\mathbf{x}^\top A \mathbf{x} > 0$ for every non-zero $\mathbf{x} \in \mathbb{R}^3$.

Solution.

- (a) Step 1: pivot 4, multipliers $\ell_{21} = -1$, $\ell_{31} = 2$. Row 2 becomes $(-4, 5, -11) +$

$(4, -4, 8) = (0, 1, -3)$, row 3 becomes $(8, -11, 27) - 2(4, -4, 8) = (0, -3, 11)$. Step 2: pivot 1, multiplier $\ell_{32} = -3$; row 3 becomes $(0, -3, 11) + 3(0, 1, -3) = (0, 0, 2)$. Hence

$$L = \begin{pmatrix} 1 & 0 & 0 \\ -1 & 1 & 0 \\ 2 & -3 & 1 \end{pmatrix}, \quad U = \begin{pmatrix} 4 & -4 & 8 \\ 0 & 1 & -3 \\ 0 & 0 & 2 \end{pmatrix},$$

and $A = LU$ by [Theorem 2.6.4](#). The diagonal 4, 1, 2 of U is non-zero, so U is invertible ([Lemma 2.6.2](#)), L is invertible ([Lemma 2.6.1](#)), and $A = LU$ is invertible. A is symmetric, so by [Theorem 2.6.8](#), with $D = \text{diag}(4, 1, 2)$,

$$A = LDL^T, \quad DL^T = \begin{pmatrix} 4 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 2 \end{pmatrix} \begin{pmatrix} 1 & -1 & 2 \\ 0 & 1 & -3 \\ 0 & 0 & 1 \end{pmatrix} = \begin{pmatrix} 4 & -4 & 8 \\ 0 & 1 & -3 \\ 0 & 0 & 2 \end{pmatrix} = U,$$

which confirms the theorem directly.

- (b) Let $\mathbf{x} \neq \mathbf{0}$ and put $\mathbf{y} = L^T \mathbf{x}$. By [Lemma 2.6.1](#), L is invertible, so L^T is invertible by Theorem 0.9.7 (4), so $\mathbf{y} \neq \mathbf{0}$ (otherwise $\mathbf{x} = (L^T)^{-1} \mathbf{y} = \mathbf{0}$). By Theorem 0.9.5, $\mathbf{y}^T = \mathbf{x}^T L$, so

$$\mathbf{x}^T A \mathbf{x} = \mathbf{x}^T L D L^T \mathbf{x} = \mathbf{y}^T D \mathbf{y} = 4y_1^2 + y_2^2 + 2y_3^2,$$

reading the 1×1 matrix as a real number. Each term is ≥ 0 , and at least one $y_i \neq 0$, which makes its term strictly positive. Hence $\mathbf{x}^T A \mathbf{x} > 0$.

Exercise C3: Counting matrices without an LU factorization.

- (a) Let $A = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \in M_2(F)$. Prove that A has an LU factorization if and only if $a \neq 0$ or $c = 0$.
- (b) Let q be a prime. How many matrices in $M_2(\mathbb{F}_q)$ have **no** LU factorization? List them for $q = 2$.

Solution.

- (a) $(\Rightarrow)$ Suppose $A = LU$ with $L = \begin{pmatrix} 1 & 0 \\ \ell & 1 \end{pmatrix}$ and $U = \begin{pmatrix} u & v \\ 0 & w \end{pmatrix}$. As in [Example \(An invertible matrix with no LU factorization\)](#), $LU = \begin{pmatrix} u & v \\ \ell u & \ell v + w \end{pmatrix}$, so $a = u$ and $c = \ell u = \ell a$. If $a = 0$, then $c = 0$. Hence $a \neq 0$ or $c = 0$.
- $(\Leftarrow)$ If $a \neq 0$, elimination without row swaps succeeds with multiplier $\ell = c/a$, and indeed

$$\begin{pmatrix} 1 & 0 \\ c/a & 1 \end{pmatrix} \begin{pmatrix} a & b \\ 0 & d - cb/a \end{pmatrix} = \begin{pmatrix} a & b \\ c & cb/a + d - cb/a \end{pmatrix} = A.$$

If $c = 0$, then A is upper triangular and $A = IA$.

- (b) By (a), A has no LU factorization exactly when $a = 0$ and $c \neq 0$. There are $q - 1$ choices of c , and q choices each of b and d , so there are $(q - 1)q^2$ such matrices. For $q = 2$ this gives 4 of the 16 matrices: $a = 0$, $c = 1$, and $b, d \in \{0, 1\}$, that is,

$$\begin{pmatrix} 0 & 0 \\ 1 & 0 \end{pmatrix}, \quad \begin{pmatrix} 0 & 0 \\ 1 & 1 \end{pmatrix}, \quad \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}, \quad \begin{pmatrix} 0 & 1 \\ 1 & 1 \end{pmatrix}.$$